

# The High-Redshift Accretion Atlas

Reproducible Catalogue and Complete Visual Growth Atlas, v3

Nicole Jiang

27 August 2026

## Abstract

We present the v3 provenance-aware atlas of JWST-identified accreting black holes at  $z \geq 4$ . The catalogue contains 142 literature measurements representing 133 physical objects in 132 host systems. Its science layer uses 119 measurements and 112 preferred physical-object rows with canonical numeric black-hole masses; 21 catalogue objects without such masses remain explicit exclusions. The atlas keeps heterogeneous evidence classes and mass methods separate, retains alternate measurements, and treats global ranks as navigation rather than demographic inference. A full-source audit corrects the Scholtz et al. JADES narrow/high-ionization subset to 21 objects at  $z \geq 4$  and documents the paper/table count discrepancy while retaining all source history. Under the reference assumptions  $z_{\text{seed}} = 30$ ,  $\epsilon = 0.1$ , and no merger boost, GN-38509 has the highest point-estimate growth pressure. These diagnostics measure pressure under stated assumptions; they do not identify a unique seed or accretion history.

## 1 Purpose and interpretation boundary

The atlas asks which JWST-identified objects place the strongest pressure on simple black-hole growth histories, and which measurements most affect that conclusion. It does not estimate a population mass function from the pooled catalogue. The input literature spans broad-line AGN selected by several JWST surveys [2, 3, 4, 5, 6, 7], host-selected broad-line candidates [9], an X-ray evidence history [10, 12], and JADES narrow/high-ionization candidates [13]. Selection functions, evidence strength, and mass estimators are therefore retained as analysis strata.

## 2 Catalogue construction and provenance

Every measurement has a stable measurement ID, physical-object ID, and host system ID. Source-native values, adopted canonical values, uncertainty kinds, method tags, URLs, available archive checksums, and identity decisions are preserved. Missing historical provenance is explicit: the 23 frozen-v1 Juodžbalis/JADES rows predate archive hashing and retain a documented unavailable checksum and extraction-date limitation rather than a fabricated value. Auxiliary coordinate and context inputs are attributed separately to the THRILS programme catalogue [8], the UHZ1 spectroscopy paper [11], and the versioned JADES DR3 GOODS-S prism catalogue. Alternate measurements are never silently averaged or deleted. Exactly one preferred measurement controls each object-level status in v3, while the complete history of measurement-level evidence statuses and bases remains in the object table. This resolves the case in which a weaker alternate candidate row had conservatively downgraded an otherwise secure preferred measurement.

The Scholtz et al. source audit deserves specific note. The paper reports 42 objects, while its source-native sample table contains 41 rows. Parsing that full table gives exactly 21 rows at  $z \geq 4$ . The pre-correction extraction contained 20 of them; canonical v3 includes the omitted JADES-NS-GS00099671

row ( $z = 5.936$ ) with its published host and bolometric observables, but no invented black-hole mass. The 21 candidates include one identity-linked broad-line object, so the current physical-object catalogue contains 20 narrow-line candidates from this source.

**Table 1:** v3 scale and explicit analysis subsets.

| Product or subset                                           | Rows / objects  |
|-------------------------------------------------------------|-----------------|
| Measurements / physical objects / host systems              | 142 / 133 / 132 |
| Source-local observables                                    | 183             |
| Growth-eligible measurements / objects                      | 119 / 112       |
| Primary measurements / objects                              | 112 / 105       |
| Broad-line eligible objects                                 | 112             |
| Catalogue-only physical objects                             | 21              |
| Alternate-measurement sensitivity pairs                     | 7               |
| Follow-up rows / ranked rows                                | 133 / 112       |
| Source-family caveat rows                                   | 9               |
| Per-object $f_{\text{Edd}}$ -mass/seed-redshift-mass panels | 266             |

A supplementary [full-assumption growth-track companion](#) preserves the historical v1 visual grid: three seed masses crossed with  $f_{\text{Edd}} = 0.3, 1, 2$ , four constant radiative efficiencies, and two merger boosts, for 72 reference curves in total. Seed mass is encoded by color,  $f_{\text{Edd}}$  by line style, efficiency by line width, and merger boost by opacity; numerical points retain the companion’s original blue palette. In the no-mass panel, narrow-line candidates are blue and the X-ray candidate is gray. This denser view is an assumption comparison, not a fit or a demographic model.

### 3 Growth model and uncertainty

The reference model follows exponential radiatively efficient growth [1],

$$M_{\text{BH}}(t_{\text{obs}}) = M_{\text{seed}} \exp \left[ f_{\text{Edd}} \left( \frac{1 - \epsilon}{\epsilon} \right) \frac{t(z_{\text{obs}}) - t(z_{\text{seed}})}{t_{\text{Edd}}} \right] B_{\text{merge}}, \quad (1)$$

with  $t_{\text{Edd}} \simeq 0.45$  Gyr. The main tables invert this relation to report required lifetime-average  $f_{\text{Edd}}$  for fixed seed masses and required  $M_{\text{seed}}$  for fixed  $f_{\text{Edd}}$ . The baseline uses  $z_{\text{seed}} = 30$ ,  $\epsilon = 0.1$ , and  $B_{\text{merge}} = 1$ .


Reported asymmetric statistical uncertainties in  $\log_{10} M_{\text{BH}}$  are propagated with 10,000 deterministic, per-object split-side draws. Method-scale systematic envelopes remain separate from the statistical sampling. A reported current Eddington ratio is not substituted for the lifetime-average growth rate.

### 4 Current results


GN-38509 has the highest point-estimate reference growth pressure, requiring  $f_{\text{Edd}} = 1.065$  for a  $10^2 M_{\odot}$  seed. RUBIES-EGS-49140 leads the uncertainty-aware navigation view, with a median required  $f_{\text{Edd}} = 1.016$ . These results express tension with that reference history, not a claim of literally continuous super-Eddington accretion.

Within-class and within-class-mass-method ranks are the supported comparison scopes. All outputs explicitly forbid pooled demographic inference without a selection-function and completeness model. The alternate-measurement table contains 7 comparisons; it shows how retained measurement choice changes  $\log_{10} M_{\text{BH}}$  and required  $f_{\text{Edd}}$  while leaving the canonical measurement choice auditable. The follow-up


matrix retains all 133 objects, ranks the 112 growth-eligible objects, and leaves the other 21 visibly unranked. A separate 9-row caveat table gives one compact provenance and interpretation summary for every admitted source family.




**Figure 1:** Current catalogue landscape. Left: all 112 growth-eligible physical objects in black-hole-mass–redshift space. Right: growth-product coverage for all 133 objects. The 21 catalogue-only objects remain available for provenance and evidence analyses but are not assigned numeric growth products.




**Figure 2:** Complete v3 growth-track view. All 112 objects with supported canonical black-hole masses are plotted against reference tracks. The lower panel retains the redshift locations of all 21 catalogue-only objects for which a numerical growth history cannot be inferred. Thus every physical object is visible without assigning unsupported masses. The common reversed redshift axis spans  $z = 10$  to  $z = 3$ , centering the observed sample while retaining the original figure dimensions and margins. Broad-line AGN are purple.




**Figure 3:** Descriptive growth-pressure distributions for the JWST broad-line AGN sample. Boxes and points summarize the canonical object-level table; they are not completeness-corrected demographic estimates.



**Figure 4:** Compatibility summary using every one of the 112 growth-computable objects. Panels retain object class and seed family while spanning spin, merger boost, and average accretion rate. Percentages are descriptive within-class fractions, not pooled demographic estimates. The complete object-by-object matrix, including all 21 unavailable objects, is included in the supplement.



**Figure 5:** Complete Monte Carlo uncertainty summary. Every one of the 112 numerical posteriors is plotted; point size encodes the 16th–84th percentile width. The adjacent accounting panel retains all 21 catalogue-only objects. The full supplement atlas gives every object label and interval.



**Figure 6:** Sensitivity to 7 retained alternate measurements. Numbered points are keyed to stable object and measurement IDs, avoiding ambiguous overlapping labels. The plot quantifies measurement-choice sensitivity rather than treating alternates as independent objects.

## 5 Reproducibility and canonical artifacts

The Python 3.12 direct environment dependencies are pinned. Canonical dataset manifests verify exact catalogue, science, and figure membership and SHA-256 hashes. In-memory reproduction compares generated tables against checked-in CSVs without rewriting them. Regression tests enforce cardinalities, source anchors, identity links, evidence aggregation, uncertainty boundaries, ranking policy, figure resolution, and every canonical gallery path. Only v1/v2/v3 products and manifests are public; the source-admission bridge required for reconstruction is isolated as internal compatibility code. Superseded release-era artifacts remain recoverable from Git history.

The final figure layer provides paper-facing summaries and a complete current gallery. Every one of the 133 objects has an  $f_{\text{Edd}}$ -mass panel and a seed-redshift-mass panel under the final gallery path. For the 21 unavailable objects these are explicit no-inference panels. Complete  $f_{\text{Edd}}$ -mass, compatibility, and Monte Carlo atlas figures preserve every object label for close zooming. Growth tracks are retained as combined all-object figures rather than individual gallery panels; the supplementary full-assumption view preserves all 72 historical v1 reference curves for direct comparison with v3.

## 6 Limitations and next scientific work

- The catalogue is heterogeneous and cannot support pooled demographics without explicit selection and completeness models.
- Virial masses are method-dependent; available method systematics are reported separately and are not full posterior reconstructions.
- Narrow-line and X-ray candidates without canonical numeric masses are valuable evidence records but cannot enter the current growth ranking.
- The reference model compresses time-variable accretion, feedback, mergers, and radiative-state changes into effective parameters.
- Identity review uses deterministic candidates and documented decisions; crowded or imprecise cases may ultimately require probabilistic crossmatching.

The next scientific milestone is no longer basic catalogue assembly. It is a selection-aware inference layer: explicit survey selection functions, completeness modeling, method-calibrated hierarchical uncertainty, and a value-of-information refinement of the existing follow-up matrix tied to which measurement would most change each object’s physical interpretation.

## 7 Conclusion

v3 is a complete, internally consistent research dataset: the current catalogue, class-aware science tables, paper figures, full gallery coverage, provenance correction, evidence policy, and automated verification all refer to the same dataset version. It provides reproducible observational triage across 133 JWST-identified objects while keeping unsupported demographic and formation-channel claims outside the scope of the data product.

## References

- [1] Dayal, P. 2024, *A&A*, 690, A182, doi:10.1051/0004-6361/202451481.



- [2] Juodžbalis, I. et al. 2026, *MNRAS*, 546, stag086, doi:10.1093/mnras/stag086.
- [3] Taylor, A. J. et al. 2025, *ApJ*, 986, 165, doi:10.3847/1538-4357/add15b.
- [4] Matthee, J. et al. 2024, *ApJ*, 963, 129, doi:10.3847/1538-4357/ad2345.
- [5] Lin, X. et al. 2024, *ApJ*, 974, 147, doi:10.3847/1538-4357/ad6565.
- [6] Harikane, Y. et al. 2023, *ApJ*, 959, 39, doi:10.3847/1538-4357/ad029e.
- [7] Davis, K. et al. 2026, arXiv:2602.23310v1.
- [8] Hutchison, T. A. et al. 2025, THRILS redshift catalogue, arXiv:2512.12509v1.
- [9] Ren, J. et al. 2025, *MNRAS*, 544, 211–233, doi:10.1093/mnras/staf1709.
- [10] Bogdán, Á. et al. 2024, *Nature Astronomy*, 8, 126, doi:10.1038/s41550-023-02111-9.
- [11] Goulding, A. D. et al. 2023, *ApJL*, 955, L24, doi:10.3847/2041-8213/acf7c5.
- [12] Zou, F. et al. 2026, UHZ1 full-data X-ray reanalysis, arXiv:2603.24893v1.
- [13] Scholtz, J. et al. 2025, *A&A*, 697, A175, doi:10.1051/0004-6361/202348804.